

## HW1, report, Dinar Sharafutdinov

Here is my report where I show the solution step by step and explain problems that I've faced

```
In [1]: %load_ext autoreload
        %autoreload 2
```

```
In [2]: from collections import defaultdict
        from copy import deepcopy
        import json
        import os
        import warnings
        import glob
        import re
        from io import StringIO, BytesIO
        from functools import partial
        from typing import Tuple

        import h5py
        import k3d
        import numpy as np
        import matplotlib.pyplot as plt
        from tqdm.notebook import tqdm, trange
```

```
In [3]: import sys
        sys.path.append('..')
```

```
In [4]: from gcv_v20211_hw1.utils.plotting import display_depth_sharpness, displ
        from gcv_v20211_hw1.utils.hdf5.dataset import Hdf5File, PreloadTypes
        from gcv_v20211_hw1.utils.sharpf_io import WholeDepthMapIO, PointPatchPr
        from gcv_v20211_hw1.utils.camera_utils.camera_pose import camera_to_disp
```

## Data preparation

```
In [5]: ground_truth_dataset = Hdf5File(
        './validation/med_res/abc_0050_00500149_54930d6f7740b03347d89a56_00
        io=WholeDepthMapIO,
        preload=PreloadTypes.LAZY,
        labels='*')

        predictions_dataset = Hdf5File(
        './validation/med_res/abc_0050_00500149_54930d6f7740b03347d89a56_00
        io=WholeDepthMapIO,
        preload=PreloadTypes.LAZY,
        labels='*')
```


```
..\gcv_v20211_hw1\utils\hdf5\dataset.py:64: UserWarning: File C:\Users\dinar\YandexDisk\skoltech-gcv-course-2021-main\skoltech-gcv-course-2021-main\HW1\gcv_v2021.1_hw1\gcv_v2021.1_hw1\validation\med_res\abc_0050_00500149_54930d6f7740b03347d89a56_000_predictions.hdf5 is not compatible with Hdf5File I/O interface <class 'gcv_v20211_hw1.utils.hdf5.io_struct.HDF5IO'>
  warnings.warn('File {} is not compatible with Hdf5File I/O interface {}'.format(
  {})).format(
```

```
In [6]: def distances_nobg(distances, image):
        # here we take distances for image pixels
        distances_ = np.zeros_like(distances)
```

```
distances_[np.nonzero(image)] = distances[np.nonzero(image)]
return distances_
```


```
In [7]: plt.imshow(ground_truth_dataset[0]['image'])
```

```
Out[7]: <matplotlib.image.AxesImage at 0x2b4e549d310>
```



```
In [8]: display_depth_sharpness(
    depth_images=[
        camera_to_display(ground_truth_dataset[0]['image']), # camera_to
        camera_to_display(ground_truth_dataset[1]['image']),
        camera_to_display(ground_truth_dataset[2]['image']),
    ],
    sharpness_images=[
        camera_to_display(
            distances_nobg(predictions_dataset[0]['distances'],
                           ground_truth_dataset[0]['image'])),
        camera_to_display(
            distances_nobg(predictions_dataset[1]['distances'],
                           ground_truth_dataset[1]['image'])),
        camera_to_display(
            distances_nobg(predictions_dataset[2]['distances'],
                           ground_truth_dataset[2]['image'])),
    ],
    axes_size=(8, 8),
    ncols=3)
```

```
..\gcv_v20211_hwl\utils\plotting.py:132: MatplotlibDeprecationWarning: Yo
u are modifying the state of a globally registered colormap. In future ve
rsions, you will not be able to modify a registered colormap in-place. To
remove this warning, you can make a copy of the colormap first. cmap = co
py.copy(mpl.cm.get_cmap("viridis"))
    depth_cmap.set_bad(color='black')
..\gcv_v20211_hwl\utils\plotting.py:148: MatplotlibDeprecationWarning: Yo
u are modifying the state of a globally registered colormap. In future ve
rsions, you will not be able to modify a registered colormap in-place. To
remove this warning, you can make a copy of the colormap first. cmap = co
py.copy(mpl.cm.get_cmap("coolwarm_r"))
    sharpness_cmap.set_bad(color='black')
```



```
In [9]: predictions_dataset.items['distances'].shape
```

```
Out[9]: (18, 1024, 1024)
```

```
In [10]: MED_RES = 0.05
         resolution_3d = MED_RES
         nn_set_size = 4
         distance_interpolation_threshold = 1.0
```

```
In [11]: images = [view['image'] for view in ground_truth_dataset]
         resolution_image = images[0].shape

         distances = [view.get('distances', np.ones_like(view['image'])) for view in ground_truth_dataset]
         extrinsics = [view['camera_pose'] for view in ground_truth_dataset]
         intrinsics = [dict(resolution_image=resolution_image, resolution_3d=resolution_3d) for view in ground_truth_dataset]
```

```
In [12]: i = 0

         image_i = images[i] # [h, w]
         distances_image_i = distances[i] # [h, w]
         # Kill background for nicer visuals
         distances_i = np.zeros_like(distances_image_i)
         distances_i[np.nonzero(image_i)] = distances_image_i[np.nonzero(image_i)]
```

Here I start to code a standard algorithm

## get\_view function

```
In [13]: # TODO: write your code to constrict a world-frame point cloud from a depth image
         # using known intrinsic and extrinsic camera parameters.
         # Hints: use the class `RaycastingImaging` to transform image to point cloud
         # use the class `CameraPose` to transform image to points in world frame
```

I use CameraPose and RaycastingImaging classes

```
In [14]: pose_i = CameraPose(extrinsics[i])
```

```
In [15]: from gcv_v20211_hw1.utils.camera_utils.imaging import RaycastingImaging
         imaging_i = RaycastingImaging(resolution_image, resolution_3d)
```

```
In [16]: points_i = imaging_i.image_to_points(image_i)
         points_i = pose_i.camera_to_world(points_i)
         points_i.shape
```

```
Out[16]: (33172, 3)
```

```
In [17]: def get_view(
         images,
         distances,
         extrinsics,
         intrinsics_dict,
         i):
         """A helper function to conveniently prepare view information."""
         image_i = images[i] # [h, w]
         distances_image_i = distances[i] # [h, w]
         # Kill background for nicer visuals
         distances_i = np.zeros_like(distances_image_i)
         distances_i[np.nonzero(image_i)] = distances_image_i[np.nonzero(image_i)]

         # TODO: write your code to constrict a world-frame point cloud from a depth image
         # using known intrinsic and extrinsic camera parameters.
         # Hints: use the class `RaycastingImaging` to transform image to point cloud
         # use the class `CameraPose` to transform image to points in world frame
```

```

# use the class `CameraPose` to transform image to points in world
pose_i = CameraPose(extrinsics[i]) #np.linalg.inv
imaging_i = RaycastingImaging(resolution_image, resolution_3d)
points_i = imaging_i.image_to_points(image_i)
points_i = pose_i.camera_to_world(points_i)

return image_i, distances_i, points_i, pose_i, imaging_i

```

```

In [18]: image_i, distances_i, points_i, pose_i, imaging_i = get_view(images, dis
_, distances_j, points_j, _, _ = get_view(images, distances, extrinsics,

distances_i_flattened = distances_i.flatten()
image_i_flattened = image_i.flatten()

```

```

In [20]: points_i, points_j;

```

Some visualisation for check

```

In [22]: plot = k3d.plot(name='points')
plot += k3d.points(points_j)
plot += k3d.points(points_i)
plot.display()

```

```

In [23]: reprojected_j = pose_i.world_to_camera(points_j)
reprojected_i = pose_i.world_to_camera(points_i)

```

```

In [24]: plot = k3d.plot(name='points')
plot += k3d.points(points_i)
plot += k3d.points(reprojected_j)
plot.display()

```

```

In [25]: im = RaycastingImaging(resolution_image, resolution_3d)
p = im.image_to_points(images[0])

```

```

In [26]: plot = k3d.plot(name='points')
plot += k3d.points(p)
plot += k3d.points(reprojected_j)
plot.display()

```

## Interpolation

```

In [ ]: # For each reprojected point, find K nearest points in view_i,
# that are source points/pixels to interpolate from.
# We do this using imaging_i.rays_origins because these
# define (u, v) coordinates of points_i in the pixel grid of view_i.
# TODO: your code here: use cKDTree to find k='nn_set_size' indexes of
# nearest points for each of points from `reprojected_j`

https://docs.scipy.org/doc/scipy/reference/generated/scipy.spatial.cKDTree
https://docs.scipy.org/doc/scipy/reference/generated/scipy.spatial.cKDTree

```

```

In [37]: from scipy.spatial import cKDTree

```

```

In [39]: uv_i = imaging_i.rays_origins[:, :2]
uv_i

```

```
Out[39]: array([[ 25.6 ,  25.6 ],
               [ 25.6 ,  25.55],
               [ 25.6 ,  25.5 ],
               ...,
               [-25.55, -25.45],
               [-25.55, -25.5 ],
               [-25.55, -25.55]])
```

```
In [40]: reprojected_j
```

```
Out[40]: array([[ -6.89849329,  0.91759927, 17.53170197],
               [ -6.89849329,  0.86313749, 17.51930121],
               [ -6.89849329,  0.8086757 , 17.50690044],
               ...,
               [ 6.79713646,  0.56047288,  7.83328935],
               [ 6.79713646,  0.5060111 ,  7.82088859],
               [ 6.79713646,  0.45154932,  7.80848782]])
```

```
In [41]: _, nn_indexes_in_i = cKDTree(uv_i).query(reprojected_j[:, :2], nn_set_s,
nn_indexes_in_i)
```

```
Out[41]: array([[666094, 666093, 665070, 667118],
               [666095, 666094, 665071, 667119],
               [666096, 666095, 665072, 667120],
               ...,
               [385525, 385524, 386549, 384501],
               [385526, 385525, 386550, 384502],
               [385527, 386551, 385526, 385528]], dtype=int64)
```

```
In [42]: # Create interpolation mask: True for points which
# can be stably interpolated (i.e. they have K neighbours present
# within a predefined radius).
interp_mask = np.zeros(len(reprojected_j)).astype(bool)
# Distances to be produces as output.
distances_j_interp = np.zeros(len(points_j), dtype=float)
```

```
In [43]: interp_mask.shape
```

```
Out[43]: (45929,)
```

```
In [47]: for idx, point_from_j in enumerate(reprojected_j[:,1, :]):
point_nn_indexes = nn_indexes_in_i[idx]
# Build an [n, 3] array of XYZ coordinates for each reprojected poin
# UV values from pixel grid and Z value from depth image.
# TODO: your code here: use `point_nn_indexes` found previously
# and distance values from `image_i` indexed by the same `point_nn_
point_from_j_nns = np.concatenate((uv_i[point_nn_indexes], image_i.f
```

```
In [50]: np.concatenate((uv_i[point_nn_indexes], image_i_flattened[point_nn_index
```

```
Out[50]: array([[ -6.9      ,  0.9      , 17.12951003],
               [ -6.9      ,  0.95     , 17.14089487],
               [ -6.85     ,  0.9      , 16.92967293],
               [ -6.95     ,  0.9      , 17.51037774]])
```

```
In [51]: point_from_j
```

```
Out[51]: array([ -6.89849329,  0.91759927, 17.53170197])
```

```
In [52]: point_from_j_nns
```

```
Out[52]: array([[ -6.9      ,  0.9      , 17.12951003],
               [ -6.9      ,  0.95     , 17.14089487],
```

```

        [-6.85      ,  0.9      , 16.92967293],
        [-6.95      ,  0.9      , 17.51037774]])

```

```

In [55]: # TODO: compute a flag indicating the possibility to interpolate
# by checking distance between `point_from_j` and its `point_from_j_nns`
# against the value of `distance_interpolation_threshold`
distances_to_nearest = np.sqrt(np.sum((point_from_j - point_from_j_nns)*
interp_mask[idx] = np.all(distances_to_nearest < distance_interpolation_

```

```

In [57]: from scipy import interpolate

```

```

# TODO: your code here: use `interpolate.interp2d`
# to construct a bilinear interpolator from distances predicted
# in `view_i` (i.e. `distances_i`) into the point in `view_j`.
# Use the interpolator to compute an interpolated distance value.
interpolator = interpolate.interp2d(uv_i[:, 0][point_nn_indexes], uv_i[:,
# distances_j_interp[idx] = ...

```

```

C:\Users\dinar\Anaconda3\envs\test\lib\site-packages\scipy\interpolate\_f
itpack_impl.py:977: RuntimeWarning: The required storage space exceeds th
e available storage space.
Probable causes: nxest or nyest too small or s is too small. (fp>s)
      kx,ky=1,1 nx,ny=4,4 m=4 fp=0.000668 s=0.000000
      warnings.warn(RuntimeWarning(_iermess2[ierm][0] + _mess))

```

```

In [71]: distances_i_flattened[point_nn_indexes]

```

```

Out[71]: array([0.111214 , 0.05993421, 0.15558059, 0.00355248])

```

```

In [59]: interpolator(point_from_j[0], point_from_j[1])

```

```

Out[59]: array([0.08097661])

```

(Bonus task, 2 points) Here I implemented the second method  
<https://docs.scipy.org/doc/scipy/reference/generated/scipy.interpolate.bisplrep>

```

In [87]: tck = interpolate.bisplrep(point_from_j_nns[:, 0], point_from_j_nns[:, 1]
znew = interpolate.bisplev(point_from_j[0], point_from_j[1], tck)
znew

```

```

Out[87]: 0.08097661340292042

```

High res data fused with this approach. There are two files interpolators and  
interpolators\_2. The second one with bisplrep interpolator

## Scripts

You can find bash script run\_fin.sh with running commands for all data. Also here are the  
cells for each data set

med\_res

```

In [46]: !python ../scripts/fuse_images.py -t ../validation/med_res/abc_0050_0050
^C

```

```

In [ ]: !python ../scripts/fuse_images.py -t ../validation/med_res/abc_0050_0050

```

```

In [ ]: !python ../scripts/fuse_images.py -t ../validation/med_res/abc_0050_0050

```

```

In [ ]: !python ../scripts/fuse_images.py -t ../validation/med_res/abc_0050_0050

```

```
In [ ]: !python ../scripts/fuse_images.py -t ../validation/med_res/abc_0050_0050
```

high\_res

```
In [276... !python3 ../scripts/fuse_images.py -t ../validation/high_res/abc_0050_0050
^C
```

```
In [ ]: !python3 ../scripts/fuse_images.py -t ../validation/high_res/abc_0050_0050
```

## Results

### Example with visualisation

```
In [30]: fused_gt_dataset = Hdf5File(
        ...     '../results/abc_0050_00500149_54930d6f7740b03347d89a56_000__ground_truth.h5',
        ...     io=PointPatchPredictionsIO,
        ...     preload=PreloadTypes.LAZY,
        ...     labels='*')

        fused_pred_dataset = Hdf5File(
        ...     '../results/abc_0050_00500149_54930d6f7740b03347d89a56_000__interpolated.h5',
        ...     io=PointPatchPredictionsIO,
        ...     preload=PreloadTypes.LAZY,
        ...     labels='*')
```

```
In [29]: fused_gt_dataset = Hdf5File(
        ...     '../results_high/abc_0050_00500166_5894bbd701b2bb0fc88a6978_007__ground_truth.h5',
        ...     io=PointPatchPredictionsIO,
        ...     preload=PreloadTypes.LAZY,
        ...     labels='*')

        fused_pred_dataset = Hdf5File(
        ...     '../results_high/abc_0050_00500166_5894bbd701b2bb0fc88a6978_007__interpolated.h5',
        ...     io=PointPatchPredictionsIO,
        ...     preload=PreloadTypes.LAZY,
        ...     labels='*')
```

```
In [31]: display_sharpness(
        ...     samples=fused_gt_dataset[0]['points'],
        ...     samples_distances=fused_pred_dataset[0]['distances']
        ... )
```

```
In [32]: plot = k3d.plot(name='points')
        plot += k3d.points(fused_gt_dataset[0]['points'])
        plot.display()
```

```
In [40]: names = glob.glob('../images/*.jpg')
        names
```

med\_res

```
In [42]: import matplotlib.pyplot as plt
        import matplotlib.image as mpimg

        f, axarr = plt.subplots(1, 5, figsize=(25,25))

        for i in range(5):
```

```
axarr[i].imshow(mpimg.imread(names[i]))
axarr[i].axis('off')
```



high\_res

```
In [44]: f, axarr = plt.subplots(1,2, figsize=(25,25))

for i in range(2):
    axarr[i].imshow(mpimg.imread(names[5+i]))
    axarr[i].axis('off')
```



## Conclusions and notes

As we can see the results are quite good. I don't really see the difference between the two types of interpolation but during the interpolation process indeed there was no warning as in interp2d version. Also, we see that her\_res results a more detailed which is reasonable

During the work I faced a problem with performance:

- Overall fusion process was very slow. In the beginning, it took around 40 hours on my laptop and 5 hours on collab for each data file. But then I realised that there is a problem with .flatten() operation inside the loop which I could use beforehand. It allowed increasing performance dramatically, to one hour in collab.

I couldn't upload results to the github, so here is a link to the cloud storage:

<https://yadi.sk/d/ftaT8HD5iG-EzQ> Link to the github repository:

[https://github.com/dinarkino/geometric\\_computer\\_vision\\_hw1](https://github.com/dinarkino/geometric_computer_vision_hw1)